

Project 1

Part 1 – Regression Analysis on Health Dataset

The health dataset consists of 5 variables (4 predictor variables and 1 response variable) and 53 observations. Below is the data description for the dataset –

X1 = death rate per 1000 residents

X2 = doctor availability per 100,000 residents

X3 = hospital availability per 100,000 residents

X4 = annual per capita income in thousands of dollars

X5 = population density people per square mile

Out of the above variables, X1 is the response variable and X2 to X5 are the predictor variables.

Subset and Model Selection

We use validation set approach and Cross-Validation for selecting the best subset and the model with least error. We see that the two variable model gives us the least MSE as **3.717631**

```
> val.errors
[1] 3.802442 3.717631 3.838414 4.408887
> which.min(val.errors)
[1] 2
```

We select the two variables as X2 and X5 as below –

```
> summary(regfit.best)
subset selection object
Call: regsubsets.formula(y ~ ., data = my_data[train, ], nvmax = 4)
4 variables (and intercept)
Forced in Forced out
X2      FALSE      FALSE
X3      FALSE      FALSE
X4      FALSE      FALSE
X5      FALSE      FALSE
1 subsets of each size up to 4
Selection Algorithm: exhaustive
      X2 X3 X4 X5
1 ( 1 ) "" "" "" ""
2 ( 1 ) "" "" "" ""
3 ( 1 ) "" "" "" ""
4 ( 1 ) "" "" "" ""
```

Now, when we use Cross-Validation method to select the best subset, we again see that it selects the two variable model to produce the least MSE as **2.618851** but selects X3 and X5 as the two variables.

```

> summary(reg.best)
Subset selection object
Call: regsubsets.formula(Y ~ ., data = my_data, nvmax = 2)
4 Variables (and intercept)
    Forced in Forced out
X2      FALSE      FALSE
X3      FALSE      FALSE
X4      FALSE      FALSE
X5      FALSE      FALSE
1 subsets of each size up to 2
Selection Algorithm: exhaustive
      x2 x3 x4 x5
1 ( 1 ) " " " " " "
2 ( 1 ) " " " " " "

```

Regression Analysis

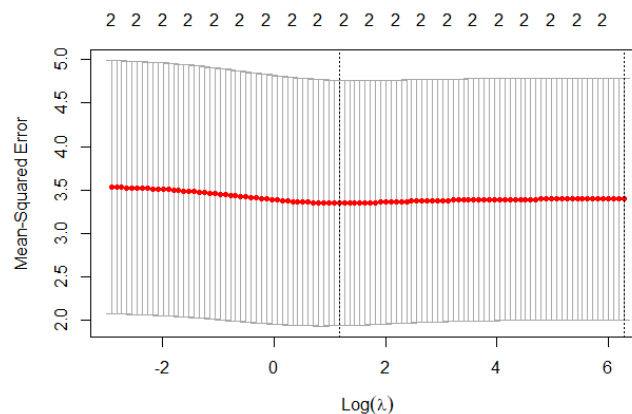
Ridge Regression

Using Cross-Validation, we find the best λ value as **3.217094** and test MSE at best λ value is **2.089405**

```

> bestlam <- cv.out$lambda.min
> bestlam
[1] 3.217094
> # we got the best lambda value as 3.217094
> # Let us now find out the test MSE at the best lambda value
> ridge.pred <- predict(ridge.mod, s = bestlam,
+                       newx = x[test, ])
> mean((ridge.pred - y.test)^2)
[1] 2.089405

```



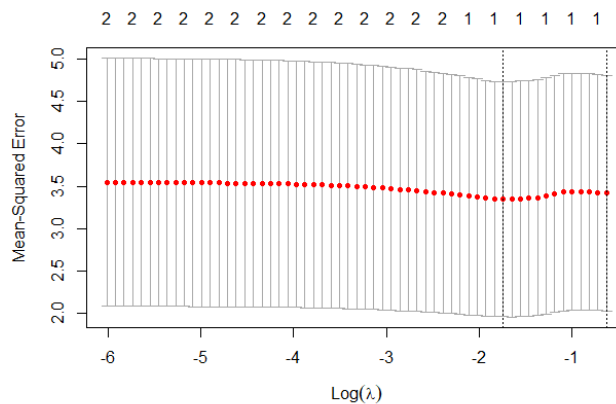
Lasso Regression

Using Cross-Validation, we find the best λ value as **0.1757266** and test MSE at best λ value is **2.103261**

```

> bestlam <- cv.out$lambda.min
> bestlam
[1] 0.1757266
> # we got the best lambda value as 0.1757266
> lasso.pred <- predict(lasso.mod, s = bestlam
+                       newx = x[test, ])
> mean((lasso.pred - y.test)^2)
[1] 2.103261
>

```



Principal Component Regression

The test MSE obtained for the dataset using PCR is **2.111575**

```
> validationplot(pcr.fit , val.type = "MSEP")
> pcr.pred <- predict(pcr.fit , x[test , ], ncomp = 2)
> mean (( pcr.pred - y.test)^2)
[1] 2.111575
```

Partial Least Squares Regression

The test MSE obtained using Partial Least Squares Regression is **2.16961**

```
> validationplot(pls.fit , val.type = "MSEP")
> pls.pred <- predict(pls.fit , x[test , ], ncomp = 1)
> mean (( pls.pred - y.test)^2)
[1] 2.16961
```

Cross – Validation

Leave One Out Cross Validation (LOOCV)

We obtained test MSE as **2.69** by using LOOCV

```
> glm.fit <- glm(Y ~ X3 + X5 , data = my_data)
> cv.err <- cv.glm(my_data , glm.fit)
> cv.err$delta
[1] 2.693229 2.690569
```

5-Fold Cross Validation

By using 5-Fold CV, we observe that the test MSE obtained is **2.539901**

```
> cv.error.5
[1] 2.759823 3.034933 2.904664 2.539901 2.647796
> which.min(cv.error.5)
[1] 4
```

10-Fold Cross Validation

By using 10-Fold CV, we observe that the test MSE obtained is **2.626308**, which is little higher than that of 5-Fold CV.

```
> cv.error.10
[1] 2.785033 2.626308 2.645378 2.773758 2.796063 2.803629 2.669285 2.767761
[9] 2.725070 2.627349
> which.min(cv.error.10)
[1] 2
```

Part 2 - Combined cycle power plant dataset

The dataset contains 9568 data points collected from a Combined Cycle Power Plant over 6 years (2006-2011), when the power plant was set to work with full load. Features consist of hourly average ambient variables Temperature (T), Ambient Pressure (AP), Relative Humidity (RH) and Exhaust Vacuum (V) to predict the net hourly electrical energy output (PE) of the plant.

A combined cycle power plant (CCPP) is composed of gas turbines (GT), steam turbines (ST) and heat recovery steam generators. In a CCPP, the electricity is generated by gas and steam turbines, which are combined in one cycle, and is transferred from one turbine to another. While the Vacuum is collected from and has effect on the Steam Turbine, the other three of the ambient variables effect the GT performance.

Attribute Information:

Features consist of hourly average ambient variables

- Temperature (T) in the range 1.81°C and 37.11°C,
- Ambient Pressure (AP) in the range 992.89-1033.30 millibar,
- Relative Humidity (RH) in the range 25.56% to 100.16%
- Exhaust Vacuum (V) in the range 25.36-81.56 cm Hg
- Net hourly electrical energy output (EP) 420.26-495.76 MW

The averages are taken from various sensors located around the plant that record the ambient variables every second. The variables are given without normalization.

Since, the variables are not scaled properly in the dataset, let us first apply log transform to normalize them as below –

```
> my_data <- read_excel("D:/CCPP.xlsx")
> summary(my_data)
      AT      V      AP      RH
Min.   :1.81  Min.   :25.36  Min.   : 992.9  Min.   : 25.56
1st Qu.:13.51 1st Qu.:41.74 1st Qu.:1009.1 1st Qu.: 63.33
Median :20.34 Median :52.08 Median :1012.9 Median : 74.97
Mean   :19.65 Mean   :54.31 Mean   :1013.3 Mean   : 73.31
3rd Qu.:25.72 3rd Qu.:66.54 3rd Qu.:1017.3 3rd Qu.: 84.83
Max.   :37.11 Max.   :81.56 Max.   :1033.3 Max.   :100.16

      PE
Min.   :420.3
1st Qu.:439.8
Median :451.6
Mean   :454.4
3rd Qu.:468.4
Max.   :495.8

> # Let us log transform the data, as it is mentioned in the data description
> #that the variables are not normalised
> log_tran_my_data = log(as.data.frame(my_data))
> summary(log_tran_my_data)
      AT      V      AP      RH
Min.   :0.5933  Min.   :3.233  Min.   :6.901  Min.   :3.241
1st Qu.:2.6034 1st Qu.:3.731 1st Qu.:6.917 1st Qu.:4.148
Median :3.0128 Median :3.953  Median :6.921  Median :4.317
Mean   :2.8881 Mean   :3.967  Mean   :6.921  Mean   :4.272
3rd Qu.:3.2473 3rd Qu.:4.198 3rd Qu.:6.925 3rd Qu.:4.441
Max.   :3.6139 Max.   :4.401  Max.   :6.941  Max.   :4.607

      PE
Min.   :6.041
1st Qu.:6.086
Median :6.113
Mean   :6.118
3rd Qu.:6.149
Max.   :6.206
```

Subset and Model Selection

We use validation set approach and Cross-Validation for selecting the best subset and the model with least error. We see that the four variable model gives us the least MSE as **0.0001060837**

```
> val.errors
[1] 0.0001735037 0.0001149345 0.0001069582 0.0001060837
> which.min(val.errors)
[1] 4
```

The cross validation tells us that the 4 variable model has the least MSE. Hence, we select that and apply best subset selection on that model and get below coefficients.

```
> mean.cv.errors <- apply(cv.errors , 2, mean)
> mean.cv.errors
      1      2      3      4
0.0001829516 0.0001182675 0.0001102586 0.0001090167

> summary(reg.best)
Subset selection object
Call: regsubsets.formula(PE ~ ., data = log_tran_my_data, nvmax = 4)
4 variables (and intercept)
  Forced in Forced out
AT      FALSE      FALSE
V       FALSE      FALSE
AP      FALSE      FALSE
RH      FALSE      FALSE
1 subsets of each size up to 4
Selection Algorithm: exhaustive

  1 ( 1 ) "x" " " " " " "
  2 ( 1 ) "x" "x" " " " "
  3 ( 1 ) "x" "x" "x" " "
  4 ( 1 ) "x" "x" "x" "x"

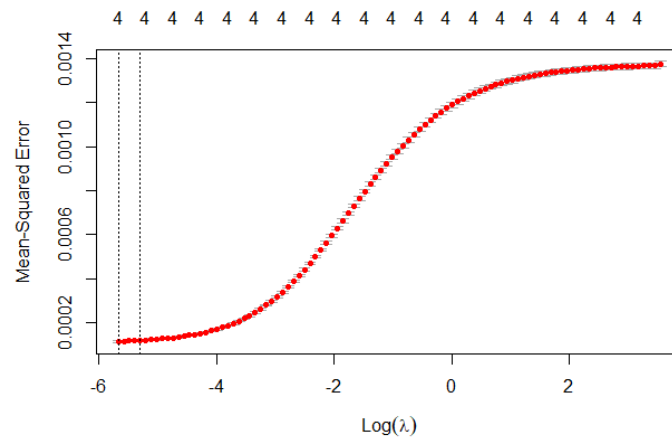
> # We see that the 4 variable model gives least MSE as 0.0001090167
> coef(reg.best , 4)
(Intercept)          AT          V          AP          RH
  2.944019024 -0.051532843 -0.055654478  0.515754969 -0.006015696
```

Regression Analysis

Ridge Regression

Using Cross-Validation, we find the best λ value as **0.003453709** and test MSE at best λ value is **0.0001361383**

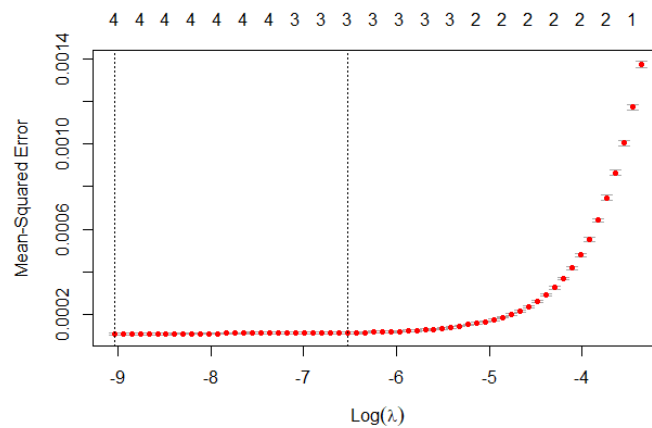
```
> bestlam <- cv.out$lambda.min
> bestlam
[1] 0.003453709
> # We got the best lambda value as 0.003453709
> # Let us now find out the test MSE at the best lambda value
> ridge.pred <- predict(ridge.mod , s = bestlam ,
+                       newx = x[test , ])
> mean (( ridge.pred - y.test)^2)
[1] 0.0001361383
```



Lasso Regression

Using Cross-Validation, we find the best λ value as **0.0001184784** and test MSE at best λ value is **0.0002372278**

```
> bestlam <- cv.out$lambda.min
> bestlam
[1] 0.0001184784
> # we got the best lambda value as 0.0001184784
> lasso.pred <- predict(lasso.mod , s = bestlam ,
+                       newx = x[test , ])
> mean (( lasso.pred - y.test)^2)
[1] 0.0002372278
```



Principal Component Regression

The test MSE obtained using PCR is **0.0001082877**

```

> summary(pcr.fit)
Data:   X dimension: 4784 4
        Y dimension: 4784 1
Fit method: svdpc
Number of components considered: 4

VALIDATION: RMSEP
Cross-validated using 10 random segments.
      (Intercept) 1 comps 2 comps 3 comps 4 comps
CV      0.03708 0.01412 0.01378 0.01125 0.01048
adjcv    0.03708 0.01412 0.01378 0.01125 0.01048

TRAINING: % variance explained
      1 comps 2 comps 3 comps 4 comps
X      59.47  81.92  95.91 100.00
PE     85.51  86.21  90.81  92.03
> validationplot(pcr.fit, val.type = "MSEP")
> pcr.pred <- predict(pcr.fit, x[test, ], ncomp = 4)
> mean((pcr.pred - y.test)^2)
[1] 0.0001082877

```

Partial Least Squares Regression

The test MSE obtained using PLS is **0.0001082877**, which is identical to that of PCR.

```

> summary(pls.fit)
Data:   X dimension: 4784 4
        Y dimension: 4784 1
Fit method: kernelpls
Number of components considered: 4

VALIDATION: RMSEP
Cross-validated using 10 random segments.
      (Intercept) 1 comps 2 comps 3 comps 4 comps
CV      0.03708 0.01284 0.01088 0.01065 0.01048
adjcv    0.03708 0.01284 0.01088 0.01065 0.01048

TRAINING: % variance explained
      1 comps 2 comps 3 comps 4 comps
X      59.29  74.9  92.29 100.00
PE     88.02  91.4  91.76  92.03
> validationplot(pls.fit, val.type = "MSEP")
> pls.pred <- predict(pls.fit, x[test, ], ncomp = 4)
> mean((pls.pred - y.test)^2)
[1] 0.0001082877

```

Cross – Validation

Leave One Out Cross Validation (LOOCV)

The test MSE obtained by using LOOCV is **0.0001090107**

```

> glm.fit <- glm(PE ~ ., data = log_tran_my_data)
> cv.err <- cv.glm(log_tran_my_data, glm.fit)
> cv.err$delta # this gives us the cross validation estimate for the test error
[1] 0.0001090108 0.0001090107
> # Linear Model

```

5-Fold Cross Validation on first degree linear model

By using 5-Fold CV on the linear model, we obtain test MSE as **0.0001089342**

```
> # Linear Model
> set.seed(15)
> cv.error.5 <- rep(0, 5)
> for (i in 1:5) {
+   glm.fit <- glm(PE ~ ., data = log_tran_my_data)
+   cv.error.5[i] <- cv.glm(log_tran_my_data, glm.fit, k = 5)$delta [1]
+ }
> cv.error.5
[1] 0.0001089342 0.0001090055 0.0001089940 0.0001090049 0.0001089445
> which.min(cv.error.5)
[1] 1
```

5-Fold Cross Validation on polynomial model

By using 5-Fold CV on the polynomial model, we obtain test MSE as **0.00008691**.

```
> cv.error.5
[1] 1.090185e-04 8.985705e-05 8.832488e-05 8.691674e-05 8.700053e-05
> which.min(cv.error.5)
[1] 4
```

10-Fold Cross Validation on first degree linear model

By using 10-Fold CV on linear model, we get test MSE as **0.0001089553**, which is almost identical to that of 5-fold CV test MSE.

```
> cv.error.10
[1] 0.0001089553 0.0001090394 0.0001089589 0.0001089797 0.0001090780
[6] 0.0001089611 0.0001090098 0.0001089801 0.0001089901 0.0001091422
> which.min(cv.error.10)
[1] 1
```

10-Fold Cross Validation on polynomial model

By using 10-Fold CV on polynomial model, we get test MSE as **0.00008456**

```
> cv.error.10
[1] 1.089636e-04 8.989113e-05 8.827954e-05 8.700801e-05 8.684348e-05
[6] 8.574404e-05 8.590240e-05 8.505376e-05 1.006417e-04 8.456811e-05
> which.min(cv.error.10)
[1] 10
```